

機器學習研究應用 Study for Machine Learning and Its Applications

Deep Learning for
Computer Vision

孫士韋

Shih-Wei Sun

swsun@newmedia.tnua.edu.tw

Outline

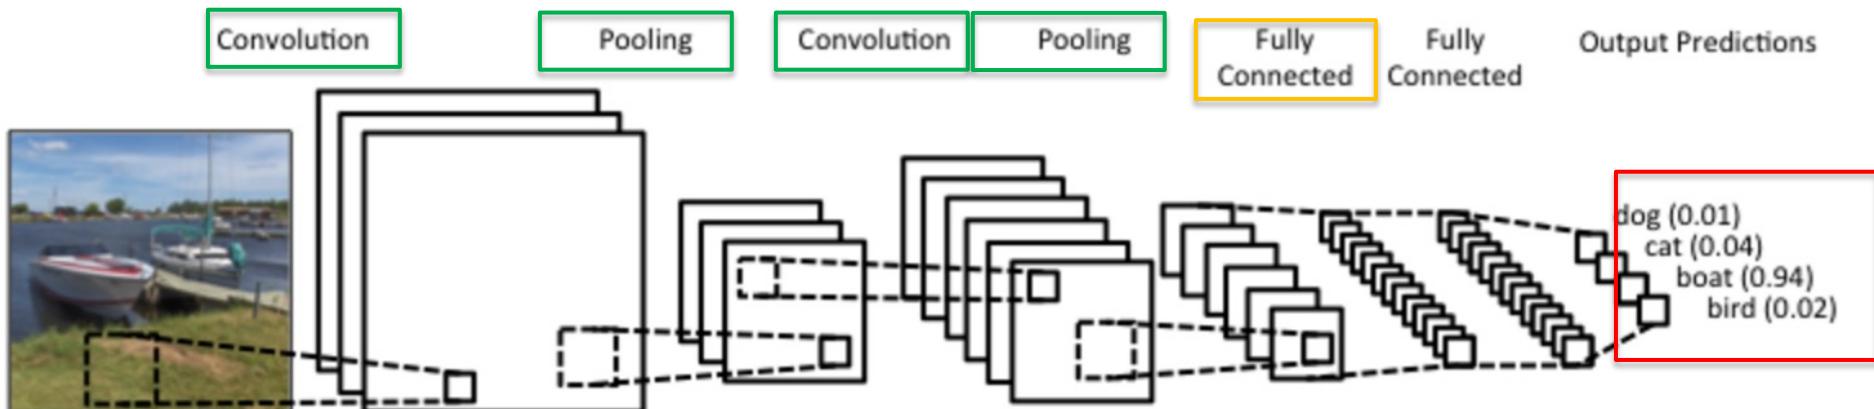
- Introduction to convnets
 - Convolution neural networks
 - Convolution operation
 - Max-pooling operation
- Training a convnet from scratch on a small dataset

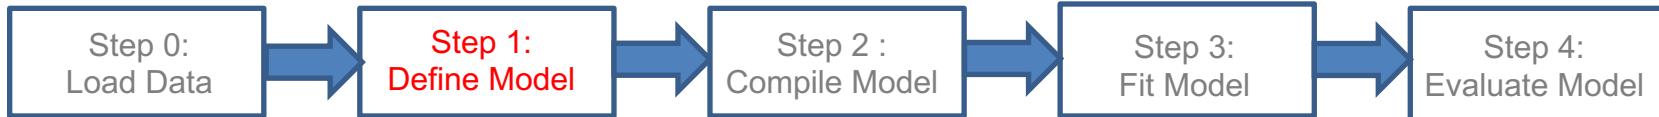
Introduction to convnets

Convolutional Neural Networks

- CNN, convnets
- Deep learning model
 - Used in computer vision applications
- Image classification problems
 - Small training datasets

Predicted
Probabilities
for each class





Introduction to convnets

- Take a practical look: a simple convnet
- Classify MNIST digits
- Instantiating a small convnet

Code:

```

1 from tensorflow.keras import layers
2 from tensorflow.keras import models
3 model = models.Sequential()      Step 1: add layers
4 model.add(layers.Conv2D(32, (3, 3), Conv2D
5     activation='relu', input_shape=(28, 28, 1)))      input image size
6 model.add(layers.MaxPooling2D((2, 2)))      MaxPooling2D
7 model.add(layers.Conv2D(64, (3, 3),
8     activation='relu'))      MaxPooling2D
9 model.add(layers.MaxPooling2D((2, 2)))
10 model.add(layers.Conv2D(64, (3, 3),
11     activation='relu'))
12
13 print(model.summary())

```

Results:

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 64)	36928
flatten (Flatten)	(None, 56)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		
None		

-2 issue:
due to convolution, border effect
(32, 64): number of channels,
Controlled by the
First argument

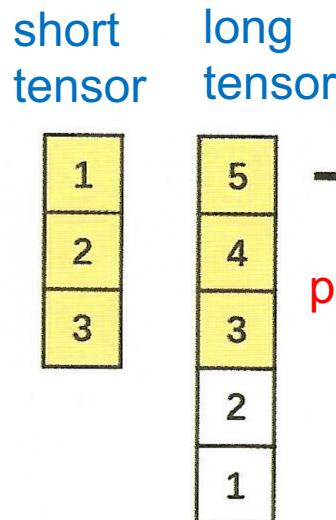
Convolution operation

- Basic concept

One kind of operations: like +, -, *, /

Convolution for two vectors:
Resulting still a vector

sliding – extracting – inner product



Slide to next



Slide to next



Boarder effect:

Original tensor: 5 elements

After convolution: 3 elements

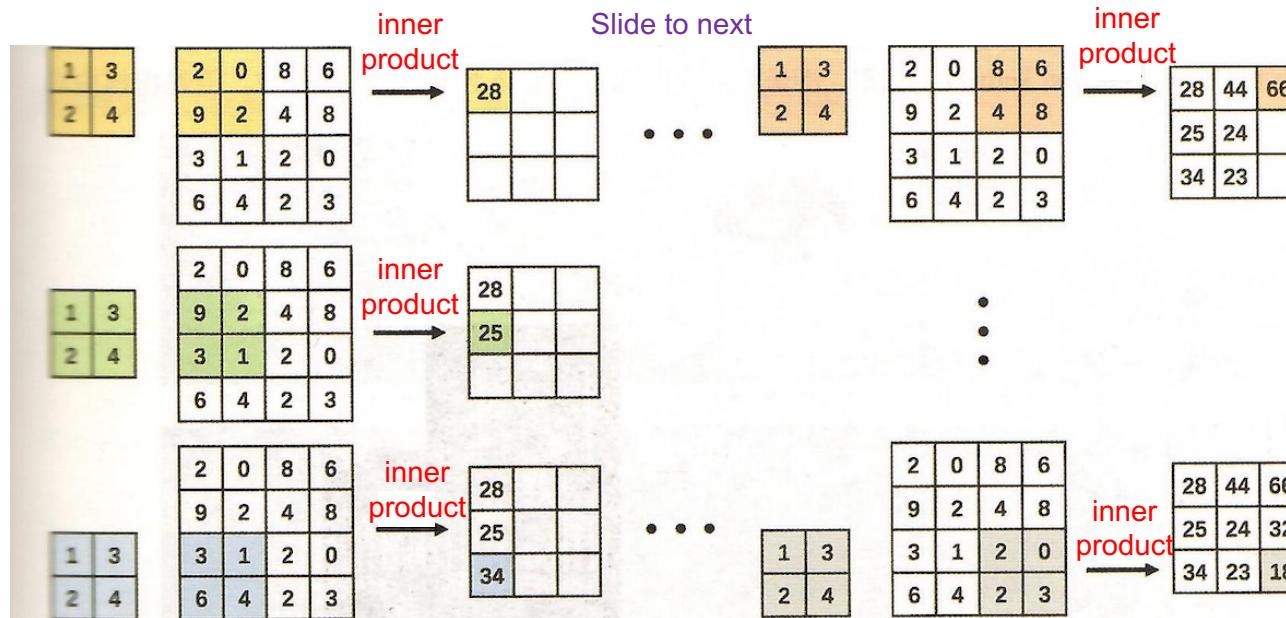
$$5 - 2 = 3$$

2D Convolution

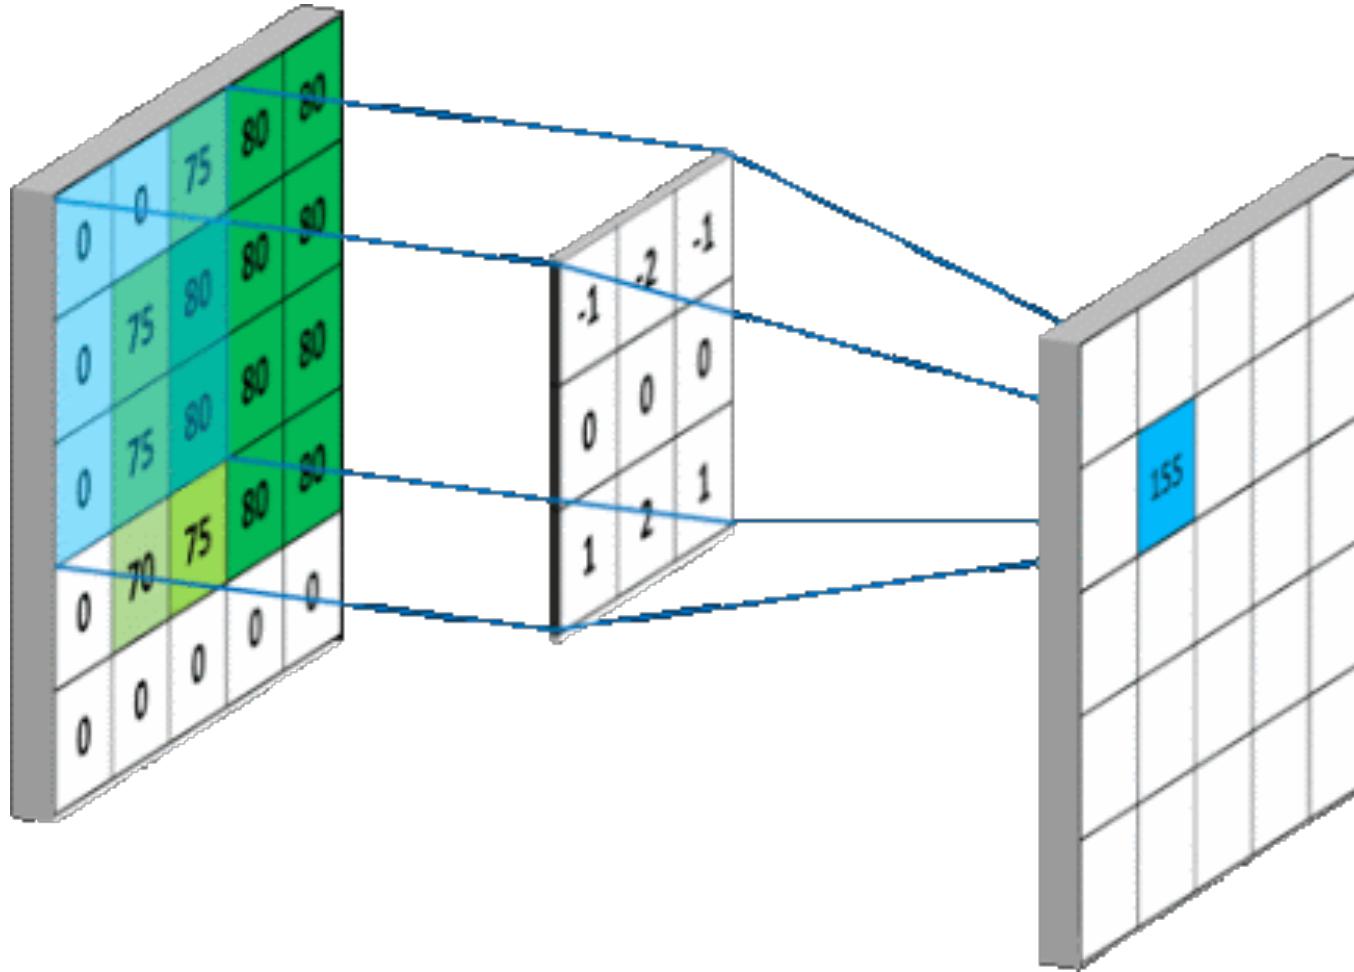
-An example of (2, 2) Conv2D

sliding – extracting – inner product

$$\begin{matrix} 1 & 3 \\ 2 & 4 \end{matrix} \bullet \begin{matrix} 0 & 3 \\ 5 & 1 \end{matrix} = 1 \times 0 + 3 \times 3 + 2 \times 5 + 4 \times 1 = 23$$



Visualization for 2D Convolution



Credit:

<https://medium.com/swlh/haar-cascade-classifiers-in-opencv-explained-visually-f608086fc42c>

The convolution operation (1/2)

- Why do we need convolution?
 - Dense layer:
 - Learn **global patterns**
 - Convolution layer:
 - Learn **Local patterns**
- Key property 1:
 - Translation invariant
 - No matter a **local pattern** moves

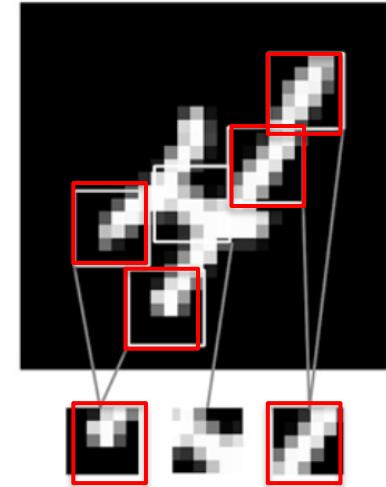
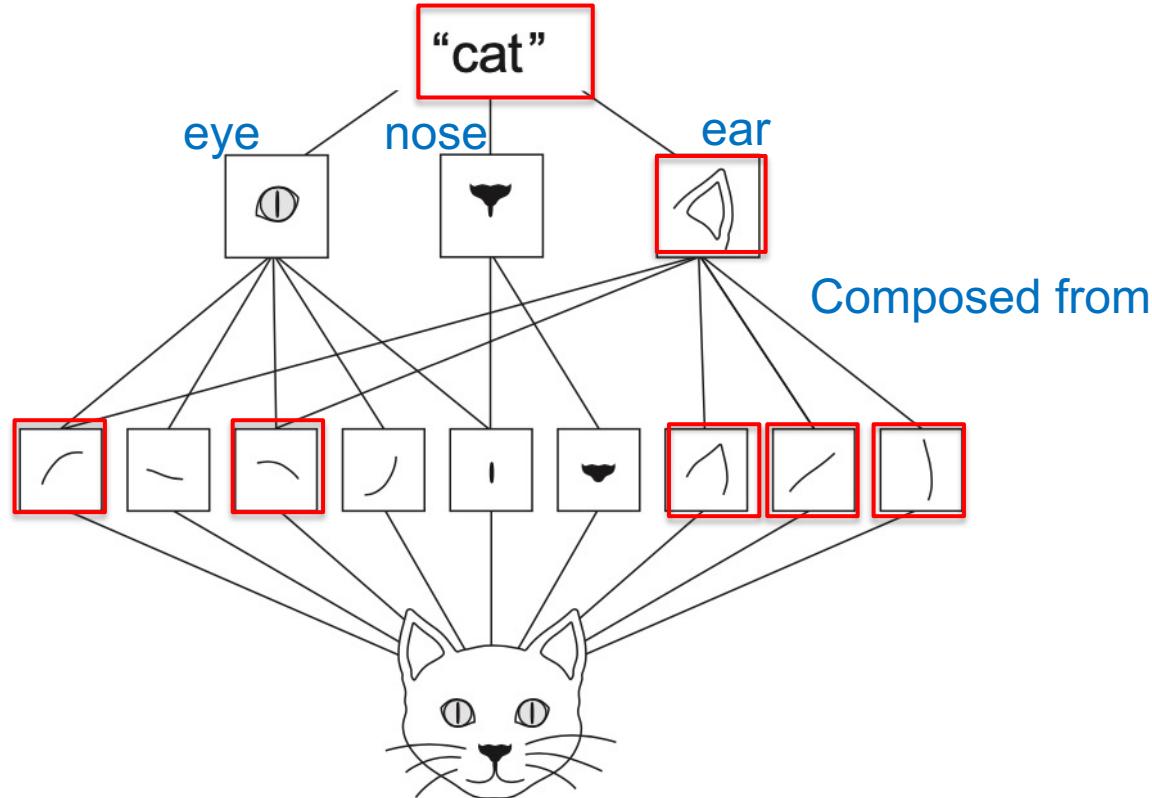


Figure 5.1 Images can be broken into local patterns such as edges, textures, and so on.

The convolution operation (2/2)

- Key property 2:
 - Learn spatial hierarchies of **local patterns**



Max-pooling operation

- Reduce the trained parameters
 - Reduce computations,
 - make the problem **solvable**

Max pooling:
- 目標downsampling
卻不會損壞識別結果

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Take the max value

Max Pooling

1		

Width, height: $\frac{1}{2}$
Area: $\frac{1}{4}$

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Take the max value

Max Pooling

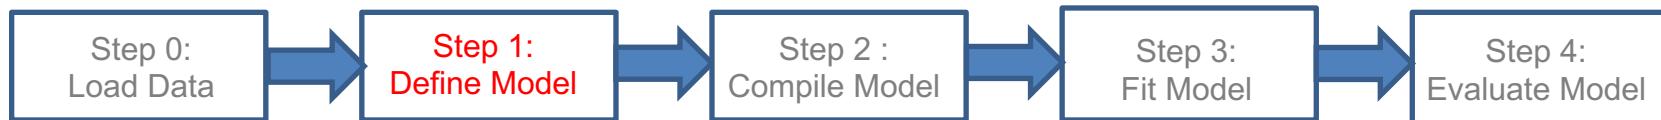
1	1	0
4	2	1
0	2	1

Feature Map

Pooled Feature Map

Feature map

- Halved
- Downsample
- MaxPooling2D
- 2×2 window



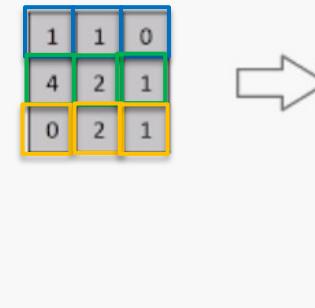
Adding a classifier to the convnet

- Feed the last output tensor
 - Into a densely connected classifier
- Flatten, dense layers

Step 1: add layers

Flatten, 3D tensor to 1D tensor

Flatten:
2D tensor -> 1D tensor



Code:

```

10 model.add(layers.Flatten())
11 model.add(layers.Dense(64, activation='relu'))
12 model.add(layers.Dense(10,
13 activation='softmax')) Add 2 dense layers
print(model.summary())

```

softmax:

multiclass, single-label problem

- Classified as 0~10 digits

Dense layer: learn global patterns

Results:

$$3 \times 3 \times 64 = 576$$

conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 10)	650
<hr/>		
Total params:	93,322	
Trainable params:	93,322	
Non-trainable params:	0	
<hr/>		
None		

2 dense layers



Training the convnet

Code:

- On MNIST images

```

15 from tensorflow.keras.datasets import mnist
16 from tensorflow.keras.utils import
     to_categorical
17
18 (train_images, train_labels), (test_images,
     test_labels) = mnist.load_data()
19
20 train_images = train_images.reshape(60000,
     28, 28, 1)
21 train_images = train_images.astype('float32')
     / 255
22 test_images = test_images.reshape(10000, 28,
     28, 1)
23 test_images = test_images.astype('float32') /
     255
24
25 train_labels = to_categorical(train_labels)
26 test_labels = to_categorical(test_labels)
27 model.compile(optimizer='rmsprop',
     loss='categorical_crossentropy',
     metrics=['accuracy'])
28 model.fit(train_images, train_labels,
     epochs=5, batch_size=64)
29 test_loss, test_acc =
     model.evaluate(test_images, test_labels)
30 print(test_acc)

```

Step 0: load data

Step 2: compile the model

Multi-class
classifier:
10 classes

Step 3: fit the model

Step 4: evaluate the model

Results:

```

Epoch 2/5
938/938 [=====] - 12s 13ms/step - loss: 0.0491 - accuracy: 0.9847
Epoch 3/5
938/938 [=====] - 12s 13ms/step - loss: 0.0329 - accuracy: 0.9896
Epoch 4/5
938/938 [=====] - 12s 13ms/step - loss: 0.0259 - accuracy: 0.9923
Epoch 5/5
938/938 [=====] - 12s 13ms/step - loss: 0.0198 - accuracy: 0.9940
2022-02-12 17:45:44.656047: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
313/313 [=====] - 2s 7ms/step - loss: 0.0406 - accuracy: 0.9887
0.9887000322341919

```

Accuracy:
98.87%

Training for 60,000 images,
 $938 \times 64 = 60,032$, close to;
for 5 epochs

Testing on 313*128
= 20,032 samples,
2 times of 10,000 images
in 2.007 second

Training need a period of time

Testing from an image
to a machine learning model,
Really fast!!

Can be applied to real-time interactive apps!

The full code

```
1 from tensorflow.keras import layers
2 from tensorflow.keras import models
3 model = models.Sequential()
4 model.add(layers.Conv2D(32, (3, 3),
5 activation='relu', input_shape=(28, 28,
6 1)))
7 model.add(layers.MaxPooling2D((2, 2)))
8 model.add(layers.Conv2D(64, (3, 3),
9 activation='relu'))
10 model.add(layers.MaxPooling2D((2, 2)))
11 model.add(layers.Conv2D(64, (3, 3),
12 activation='relu'))
13 model.add(layers.Flatten())
14 model.add(layers.Dense(64, activation='relu'))
15 model.add(layers.Dense(10,
16 activation='softmax'))
17 print(model.summary())
18
```

```
15 from tensorflow.keras.datasets import mnist
16 from tensorflow.keras.utils import
17     to_categorical
18 (train_images, train_labels), (test_images,
19     test_labels) = mnist.load_data()
20 train_images = train_images.reshape((60000,
21     28, 28, 1))
22 train_images = train_images.astype('float32') /
23     255
24 test_images = test_images.reshape((10000, 28,
25     28, 1))
26 test_images = test_images.astype('float32') /
27     255
28 train_labels = to_categorical(train_labels)
29 test_labels = to_categorical(test_labels)
30 model.compile(optimizer='rmsprop',
31                 loss='categorical_crossentropy',
32                 metrics=['accuracy'])
33 model.fit(train_images, train_labels,
34             epochs=5, batch_size=64)
```

Practice 1

- Show the accuracy

```
Epoch 2/5
938/938 [=====] - 12s 13ms/step - loss: 0.0491 - accuracy: 0.9847
Epoch 3/5
938/938 [=====] - 12s 13ms/step - loss: 0.0329 - accuracy: 0.9896
Epoch 4/5
938/938 [=====] - 12s 13ms/step - loss: 0.0259 - accuracy: 0.9923
Epoch 5/5
938/938 [=====] - 12s 13ms/step - loss: 0.0198 - accuracy: 0.9940
2022-02-12 17:45:44.656047: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
313/313 [=====] - 2s 7ms/step - loss: 0.0406 - accuracy: 0.9887
0.9887000322341919
```

Outline

- Introduction to convnets
- Training a convnet from scratch on a small dataset
 - Starting from a real example
 - Binary classification for **dogs /cats**

Training a convnet from scratch on a
small dataset

Training from a small dataset

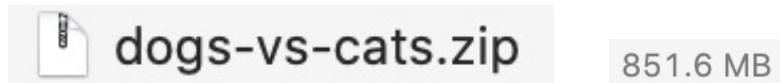
- Download the Dogs vs. Cats dataset [2013]
 - <https://www.kaggle.com/c/dogs-vs-cats/data>
 - Create a **free** user account
 - https://www.dropbox.com/s/dbjkh21m6nyxrd4/train_cats_dogs.zip?dl=0
- **25,000** images of dogs and cats
 - 12,500 dogs
 - 12,500 cats



Figure 5.8 Samples from the Dogs vs. Cats dataset. Sizes weren't modified: the samples are heterogeneous in size, appearance, and so on.

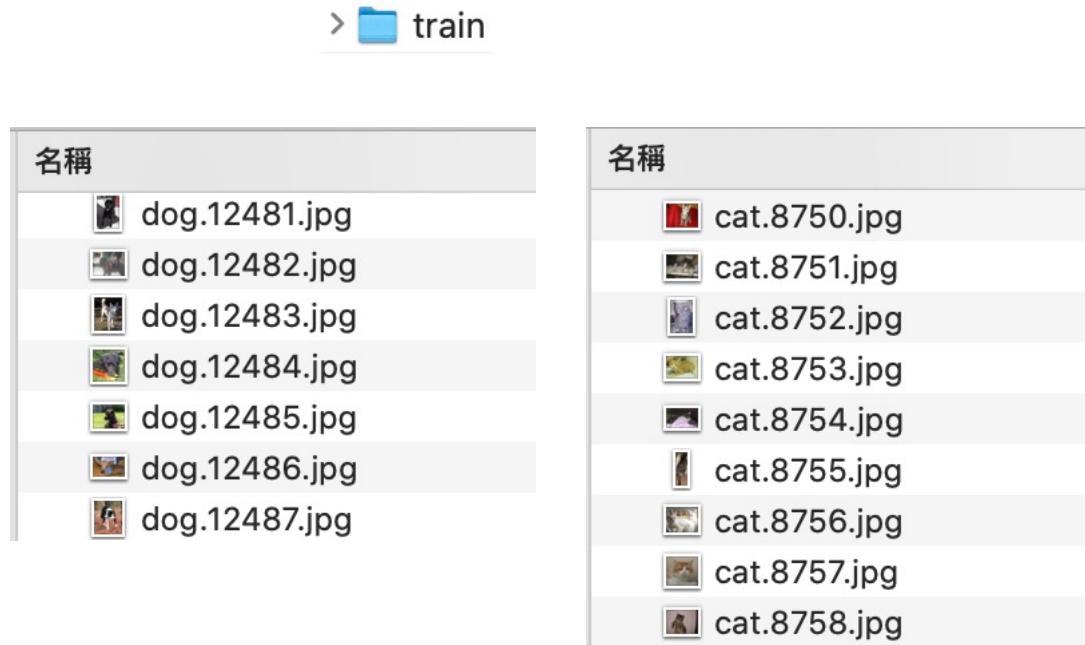
Preprocessing - Copy images

- After downloading
 - Uncompressing it
- Train your model on less than
 - 10% of the data
- To training, validation, and test directories
 - 1,000 samples/class: training
 - 500 samples/class: validation set



Uncompressing the files

- Uncompressing the zip file
- Rename the folder: train



Setup the directories

- Assigning for
 - the original file folder: images of *dogs/cats*
 - the a **smaller dataset folder**:
 - 1000 samples/training, 500 samples/class: validation

```
1 import os, shutil
2
3 # The path to the directory where the
4 # original
5 # dataset was uncompressed
6 original_dataset_dir =
7     '/Users/swsun/Downloads/train'
8
9 # The directory where we will
10 # store our smaller dataset
11 base_dir =
12     '/Users/swsun/Downloads/cats_and_dogs_s
13     mall'
14
15 os.mkdir(base_dir)    Make a new directory
```

Copy the files (1/2)

Copy 1000 cats images
to training set

```
11
12 # Directories for our training,
13 # validation and test splits
14 train_dir = os.path.join(base_dir, 'train')
15 os.mkdir(train_dir)
16 validation_dir = os.path.join(base_dir,
17     'validation')
18 os.mkdir(validation_dir)
19 test_dir = os.path.join(base_dir, 'test')
20 os.mkdir(test_dir)
```

Making folders of sets:
train/
validation/
test

```
21 # Directory with our training cat pictures
22 train_cats_dir = os.path.join(train_dir,
23     'cats')
24 os.mkdir(train_cats_dir)
25
26 # Directory with our training dog pictures
27 train_dogs_dir = os.path.join(train_dir,
28     'dogs')
29 os.mkdir(train_dogs_dir)
```

training set:
cats/dogs

```
30 validation_cats_dir =
31     os.path.join(validation_dir, 'cats')
32 os.mkdir(validation_cats_dir)
33
34 validation_dogs_dir =
35     os.path.join(validation_dir, 'dogs')
36 os.mkdir(validation_dogs_dir)
```

validation set:
cats/dogs

```
37 test_cats_dir = os.path.join(test_dir,
38     'cats')
39 os.mkdir(test_cats_dir)
40
41 test_dogs_dir = os.path.join(test_dir,
42     'dogs')
43 os.mkdir(test_dogs_dir)
```

test set:
cats/dogs

```
45 # Copy first 1000 cat images to
46     train_cats_dir
47 fnames = ['cat.{}.jpg'.format(i) for i in
48     range(1000)]
49 for fname in fnames:
50     src =
51         os.path.join(original_dataset_dir,
52             fname)
53     dst = os.path.join(train_cats_dir,
54         fname)
55     shutil.copyfile(src, dst)
```

Copy 500 cats images
to validation set

```
52 # Copy next 500 cat images to
53     validation_cats_dir
54 fnames = ['cat.{}.jpg'.format(i) for i in
55     range(1000, 1500)]
56 for fname in fnames:
57     src =
58         os.path.join(original_dataset_dir,
59             fname)
60     dst =
61         os.path.join(validation_cats_dir,
62             fname)
63     shutil.copyfile(src, dst)
```

Copy 500 cats images
to test set

```
58 # Copy next 500 cat images to test_cats_dir
59 fnames = ['cat.{}.jpg'.format(i) for i in
60     range(1500, 2000)]
61 for fname in fnames:
62     src =
63         os.path.join(original_dataset_dir,
64             fname)
65     dst = os.path.join(test_cats_dir,
66         fname)
67     shutil.copyfile(src, dst)
```

Copy 1000 dogs images
to training set

```
66 # Copy first 1000 dog images to
67     train_dogs_dir
68 fnames = ['dog.{}.jpg'.format(i) for i in
69     range(1000)]
70 for fname in fnames:
71     src =
72         os.path.join(original_dataset_dir,
73             fname)
74     dst = os.path.join(train_dogs_dir,
75         fname)
76     shutil.copyfile(src, dst)
```

Copy the files (2/2)

Copy 500 dogss images
to validation set

```
72
73 # Copy next 500 dog images to validation_dogs_dir
74 fnames = ['dog.{}.jpg'.format(i) for i in
75     range(1000, 1500)]
76 for fname in fnames:
77     src = os.path.join(original_dataset_dir,
78         fname)
79     dst = os.path.join(validation_dogs_dir,
80         fname)
81     shutil.copyfile(src, dst)
82
83
84
85
86
```

```
80 # Copy next 500 dog images to test_dogs_dir
81 fnames = ['dog.{}.jpg'.format(i) for i in
82     range(1500, 2000)]
83 for fname in fnames:
84     src = os.path.join(original_dataset_dir,
85         fname)
86     dst = os.path.join(test_dogs_dir, fname)
87     shutil.copyfile(src, dst)
```

Copy 500 dogs images
to test set

Display the number images in
Each directories

```
87 ## leave as a file for download
88 print('total training cat images:', len(os.listdir(train_cats_dir)))
89 print('total training dog images:', len(os.listdir(train_dogs_dir)))
90 print('total validation cat images:', len(os.listdir(validation_cats_dir)))
91 print('total validation dog images:', len(os.listdir(validation_dogs_dir)))
92 print('total test cat images:', len(os.listdir(test_cats_dir)))
93 print('total test dog images:', len(os.listdir(test_dogs_dir)))
```

Results:

```
total training cat images: 1000
total training dog images: 1000
total validation cat images: 500
total validation dog images: 500
total test cat images: 500
total test dog images: 500
```



Comment the codes

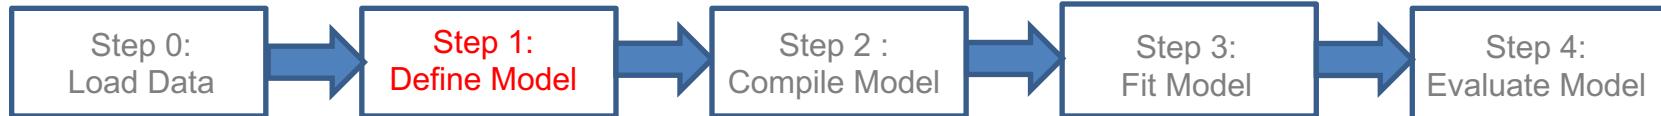
```
1 import os, shutil
2
3 # The path to the directory where the original
4 # dataset was uncompressed
5 original_dataset_dir =
6     '/Users/swsun/Downloads/train'
7
8 # The directory where we will
9 # store our smaller dataset
10 base_dir =
11     '/Users/swsun/Downloads/cats_and_dogs_small'
12 #os.mkdir(base_dir)
13
14 # Directories for our training
15 # validation and test splits
16 train_dir = os.path.join(base_dir, 'train')
17 #os.mkdir(train_dir)
18 validation_dir = os.path.join(base_dir,
19     'validation')
20 #os.mkdir(validation_dir)
21 test_dir = os.path.join(base_dir, 'test')
22 #os.mkdir(test_dir)
23
24
25 # Directory with our training cat pictures
26 train_cats_dir = os.path.join(train_dir, 'cats')
27 #os.mkdir(train_cats_dir)
28
29 # Directory with our training dog pictures
30 train_dogs_dir = os.path.join(train_dir, 'dogs')
31 #os.mkdir(train_dogs_dir)
32
33 # Directory with our validation cat pictures
34 validation_cats_dir =
35     os.path.join(validation_dir, 'cats')
36 #os.mkdir(validation_cats_dir)
37
38 # Directory with our validation dog pictures
39 validation_dogs_dir =
40     os.path.join(validation_dir, 'dogs')
41 #os.mkdir(validation_dogs_dir)
42
43 #os.mkdir(test_dogs_dir)
44 """
45 # Copy first 1000 cat images to train_cats_dir
46 fnames = ['cat.{}.jpg'.format(i) for i in
47             range(1000)]
48 for fname in fnames:
49     src = os.path.join(original_dataset_dir,
50         fname)
51     dst = os.path.join(train_cats_dir, fname)
52     shutil.copyfile(src, dst)
53
54 # Copy next 500 cat images to validation_cats_dir
55 fnames = ['cat.{}.jpg'.format(i) for i in
56             range(1000, 1500)]
57 for fname in fnames:
58     src = os.path.join(original_dataset_dir,
59         fname)
60     dst = os.path.join(validation_cats_dir,
61         fname)
62     shutil.copyfile(src, dst)
63
64 # Copy next 500 cat images to test_cats_dir
65 fnames = ['cat.{}.jpg'.format(i) for i in
66             range(1500, 2000)]
```

Comment for:
mkdir

```
32
33 # Directory with our validation dog pictures
34 validation_dogs_dir =
35     os.path.join(validation_dir, 'dogs')
36 #os.mkdir(validation_dogs_dir)
37
38 # Directory with our validation cat pictures
39 test_cats_dir = os.path.join(test_dir, 'cats')
40 #os.mkdir(test_cats_dir)
41
42 # Directory with our validation dog pictures
43 test_dogs_dir = os.path.join(test_dir, 'dogs')
44 #os.mkdir(test_dogs_dir)
45 """
46 # Copy first 1000 cat images to train_cats_dir
47 fnames = ['cat.{}.jpg'.format(i) for i in
48             range(1000)]
49 for fname in fnames:
50     src = os.path.join(original_dataset_dir,
51         fname)
52     dst = os.path.join(train_cats_dir, fname)
53     shutil.copyfile(src, dst)
54
55 # Copy next 500 cat images to validation_cats_dir
56 fnames = ['cat.{}.jpg'.format(i) for i in
57             range(1000, 1500)]
58 for fname in fnames:
59     src = os.path.join(original_dataset_dir,
60         fname)
61     dst = os.path.join(validation_cats_dir,
62         fname)
63     shutil.copyfile(src, dst)
64
65 # Copy next 500 cat images to test_cats_dir
66 fnames = ['cat.{}.jpg'.format(i) for i in
67             range(1500, 2000)]
```

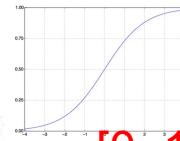
Comment for: multiple lines

```
61 for fname in fnames:
62     src = os.path.join(original_dataset_dir,
63         fname)
64     dst = os.path.join(test_cats_dir, fname)
65     shutil.copyfile(src, dst)
66
67 # Copy first 1000 dog images to train_dogs_dir
68 fnames = ['dog.{}.jpg'.format(i) for i in
69             range(1000)]
70 for fname in fnames:
71     src = os.path.join(original_dataset_dir,
72         fname)
73     dst = os.path.join(train_dogs_dir, fname)
74     shutil.copyfile(src, dst)
75
76 # Copy next 500 dog images to validation_dogs_dir
77 fnames = ['dog.{}.jpg'.format(i) for i in
78             range(1000, 1500)]
79 for fname in fnames:
80     src = os.path.join(original_dataset_dir,
81         fname)
82     dst = os.path.join(validation_dogs_dir,
83         fname)
84     shutil.copyfile(src, dst)
85
86
87 # leave as a file for download
88 print('total training cat images:', len(os.listdir(train_cats_dir)))
89 print('total training dog images:', len(os.listdir(train_dogs_dir)))
90 print('total validation cat images:', len(os.listdir(validation_cats_dir)))
91 print('total validation dog images:', len(os.listdir(validation_dogs_dir)))
92 print('total test cat images:', len(os.listdir(test_cats_dir)))
93 print('total test dog images:', len(os.listdir(test_dogs_dir)))
94 """
```



Build your network

Problem type	Last-layer activation
Binary classification	sigmoid



[0, 1]

MaxPooling2D (2,2)

Step1: add layers

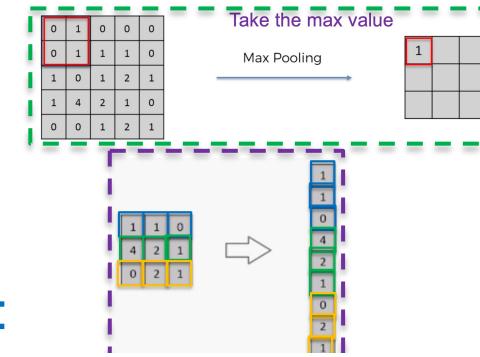
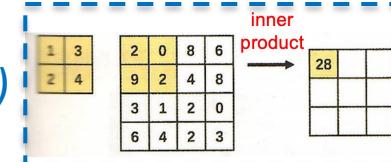
```

95
96 from tensorflow.keras import layers
97 from tensorflow.keras import models
98
99 model = models.Sequential()
100 model.add(layers.Conv2D(32, (3, 3),
101     activation='relu', input_shape=(150, 150,
102     3)))
103 model.add(layers.MaxPooling2D((2, 2)))
104 model.add(layers.Conv2D(64, (3, 3),
105     activation='relu'))
106 model.add(layers.MaxPooling2D((2, 2)))
107 model.add(layers.Conv2D(128, (3, 3),
108     activation='relu'))
109 model.add(layers.MaxPooling2D((2, 2)))
110 model.add(layers.Flatten())
111 model.add(layers.Dense(512, activation='relu'))
112 model.add(layers.Dense(1, activation='sigmoid'))
113
114 print(model.summary())

```

2 dense layers

'sigmoid' activation:
Binary classification: dog / cat



Flatten()

Results:

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		



Configuring the model

- Compile the model

```

113
114 from tensorflow.keras import optimizers
115
116 model.compile(loss='binary_crossentropy',
117                 optimizer=optimizers.RMSprop(lr=1e-4),
118                 metrics=['acc'])

```

Step 2: compile the model

Assign the loss function: binary_crossentropy

Optimizer: RMSprop, as usual

Table 4.1 Choosing the right last-layer activation and loss function for your model

Problem type	Last-layer activation	Loss function
Binary classification <i>dog / cat</i>	sigmoid	binary_crossentropy

Results:

```

flatten_1 (Flatten)      (None, 6272)
-----
dense_1 (Dense)          (None, 512)
-----
dense_2 (Dense)          (None, 1)
-----
Total params: 3,453,121
Trainable params: 3,453,121
Non-trainable params: 0
-----
```

Nothing changed



Data Preprocessing

```

118 from tensorflow.keras.preprocessing.image import
      ImageDataGenerator
119
120 # All images will be rescaled by 1./255
121 train_datagen =
      ImageDataGenerator(rescale=1./255)
122 test_datagen = ImageDataGenerator(rescale=1./255)
123
      Rescale by 1/255
124 train_generator =
      train_datagen.flow_from_directory(
      train_dir, target_size=(150, 150),
      batch_size=20, class_mode='binary')
125
126 validation_generator =
      test_datagen.flow_from_directory(
      validation_dir, target_size=(150, 150),
      batch_size=20, class_mode='binary')
127
      validation set
128 for data_batch, labels_batch in train_generator:
      print('data batch shape:', data_batch.shape) 150, 150: image tensors
129      print('labels batch shape:',
      labels_batch.shape)
130      break
131
      Display Generator results
      None
      Found 2000 images belonging to 2 classes.
      Found 1000 images belonging to 2 classes.
      data batch shape: (20, 150, 150, 3)
      labels batch shape: (20, 2)
      20. batch size

```

Preprocessing:

- Read the pictures
- Decode from JPEG to RGB
- Convert to floating point tensors
- Rescale from [0, 255] to [0,1]

ImageDataGenerator

training set

Resize to 150 x 150

Binary labels: dog/cat

validation set

Results:

Display Generator results

None

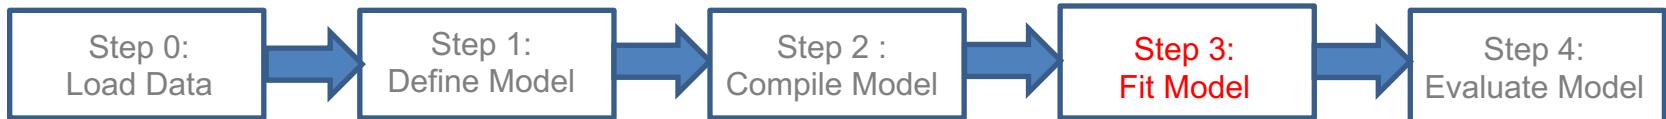
Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

data batch shape: (20, 150, 150, 3)

labels batch shape: (20, 2)

20. batch size



Fitting and saving the model

- Fit the model instead of `model.fit`
- Save the model! Pass argument:
`train_generator, validation_generator`

```

132
133 history = model.fit_generator(train_generator,
                                 steps_per_epoch=100, epochs=30,
                                 validation_data=validation_generator,
                                 validation_steps=50)
134
model
    .save
    ('/Users/swsun/Downloads/cats_and_dogs_small_1.h5')
135

```

Save the model after training

Training cost a lot of time,
Don't forget to save it!
- To be used by other interactive applications!

Results:

```

Found 2000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.
data batch shape: (20, 150, 150, 3)
labels batch shape: (20,)
Epoch 1/30
100/100 [=====] - 37s 373ms
0.5155 - val_loss: 0.6832 - val_acc: 0.5000
Epoch 2/30
100/100 [=====] - 40s 402ms
0.5850 - val_loss: 0.6617 - val_acc: 0.6370
Epoch 3/30
100/100 [=====] - 40s 402ms
0.6530 - val_loss: 0.7222 - val_acc: 0.6510
Epoch 4/30
31/100 [=====] - ETA: 22s

```



```

Epoch 27/30
100/100 [=====] - 0.9790 - val_loss: 0.0939 - val_acc: 0.7480
Epoch 28/30
100/100 [=====] - 9855 - val_loss: 1.1271 - val_acc: 0.7400
Epoch 29/30
100/100 [=====] - 0.9865 - val_loss: 0.3695 - val_acc: 0.7320
Epoch 30/30
100/100 [=====] - 0.9920 - val_loss: 0.7260 - val_acc: 0.7370

```

Practice 2

- Display the accuracy of
 - Training data
 - Validation data

```
Epoch 27/30
100/100 [=====] - 36s 359ms/step - loss: 0.0683 - acc: 0.9790 - val_loss: 0.0939 - val_acc: 0.7480
Epoch 28/30
100/100 [=====] - 322s 3s/step - loss: 0.0559 - acc: 0.9855 - val_loss: 1.1271 - val_acc: 0.7400
Epoch 29/30
100/100 [=====] - 1760s 18s/step - loss: 0.0501 - acc: 0.9865 - val_loss: 0.3695 - val_acc: 0.7320
Epoch 30/30
100/100 [=====] - 44s 437ms/step - loss: 0.0423 - acc: 0.9920 - val_loss: 0.7260 - val_acc: 0.7370
```

Plot the results

```
136 import matplotlib.pyplot as plt  
137  
138 acc = history.history['acc']  
139 val_acc = history.history['val_acc']  
140 loss = history.history['loss']  
141 val_loss = history.history['val_loss']  
142  
143 epochs = range(len(acc))  
144  
145 plt.plot(epochs, acc, 'bo', label='Training acc')  
146 plt.plot(epochs, val_acc, 'b', label='Validation  
    acc')  
147 plt.title('Training and validation accuracy')  
148 plt.legend()  
149  
150 plt.figure()  
151  
152 plt.plot(epochs, loss, 'bo', label='Training  
    loss')  
153 plt.plot(epochs, val_loss, 'b',  
    label='Validation loss')  
154 plt.title('Training and validation loss')  
155 plt.legend()  
156  
157 plt.show()
```

loss:

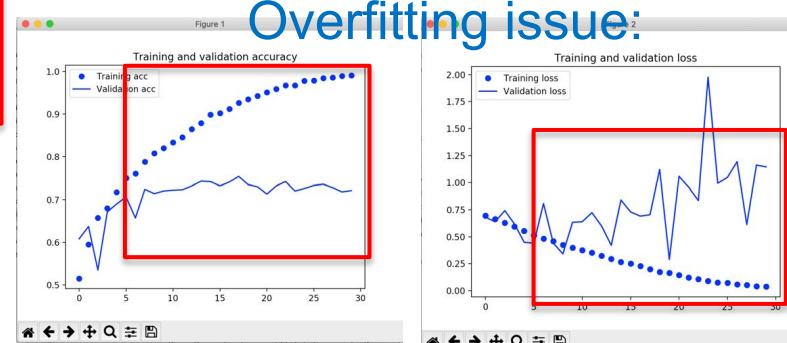
Training / validation

Accuracy:
Training / validation

Overfitting issue:

- Training accuracy – 100%
- Validation accuracy – 70%
- after 5 epochs

Too few samples: 2000 images



Improvement: Dropout,
weight decay (L2 regularization),
Reduce network size...

Full codes (1/3)

```
1 import os, shutil
2
3 # The path to the directory where the original
4 # dataset was uncompressed
5 original_dataset_dir =
6     '/Users/swsun/Downloads/train'
7
8 # The directory where we will
9 # store our smaller dataset
10 base_dir =
11     '/Users/swsun/Downloads/cats_and_dogs_small'
12 #os.mkdir(base_dir)
13
14 # Directories for our training,
15 # validation and test splits
16 train_dir = os.path.join(base_dir, 'train')
17 #os.mkdir(train_dir)
18 validation_dir = os.path.join(base_dir,
19     'validation')
20 #os.mkdir(validation_dir)
21 test_dir = os.path.join(base_dir, 'test')
22 #os.mkdir(test_dir)
23
24 # Directory with our training cat pictures
25 train_cats_dir = os.path.join(train_dir, 'cats')
26 #os.mkdir(train_cats_dir)
27
28 # Directory with our training dog pictures
29 train_dogs_dir = os.path.join(train_dir, 'dogs')
30 #os.mkdir(train_dogs_dir)
31
32 # Directory with our validation cat pictures
33 validation_cats_dir =
34     os.path.join(validation_dir, 'cats')
35 #os.mkdir(validation_cats_dir)
36
37 # Directory with our validation dog pictures
38 validation_dogs_dir =
39     os.path.join(validation_dir, 'dogs')
40 #os.mkdir(validation_dogs_dir)
41
42 # Directory with our test cat pictures
43 test_cats_dir = os.path.join(test_dir, 'cats')
44 #os.mkdir(test_cats_dir)
45
46 # Directory with our test dog pictures
47 test_dogs_dir = os.path.join(test_dir, 'dogs')
48 #os.mkdir(test_dogs_dir)
49
50 # Copy first 1000 cat images to train_cats_dir
51 fnames = ['cat.{}.jpg'.format(i) for i in
52     range(1000)]
53 for fname in fnames:
54     src = os.path.join(original_dataset_dir,
55         fname)
56     dst = os.path.join(train_cats_dir, fname)
57     shutil.copyfile(src, dst)
58
59 # Copy next 500 cat images to validation_cats_dir
60 fnames = ['cat.{}.jpg'.format(i) for i in
61     range(1000, 1500)]
62 for fname in fnames:
63     src = os.path.join(original_dataset_dir,
64         fname)
65     dst = os.path.join(validation_cats_dir, fname)
66     shutil.copyfile(src, dst)
67
68 # Copy first 1000 dog images to train_dogs_dir
69 fnames = ['dog.{}.jpg'.format(i) for i in
70     range(1000)]
71 for fname in fnames:
72     src = os.path.join(original_dataset_dir,
73         fname)
74     dst = os.path.join(train_dogs_dir, fname)
75     shutil.copyfile(src, dst)
76
77 # Copy next 500 dog images to validation_dogs_dir
78 fnames = ['dog.{}.jpg'.format(i) for i in
79     range(1000, 1500)]
80 for fname in fnames:
81     src = os.path.join(original_dataset_dir,
82         fname)
83     dst = os.path.join(validation_dogs_dir, fname)
84     shutil.copyfile(src, dst)
85
86 ## leave as a file for download
87 print('total training cat images:', len(os.listdir(train_cats_dir)))
88 print('total training dog images:', len(os.listdir(train_dogs_dir)))
89 print('total validation cat images:', len(os.listdir(validation_cats_dir)))
90 print('total validation dog images:', len(os.listdir(validation_dogs_dir)))
91 print('total test cat images:', len(os.listdir(test_cats_dir)))
92 print('total test dog images:', len(os.listdir(test_dogs_dir)))
93 print('...')
```

Full code (2/3)

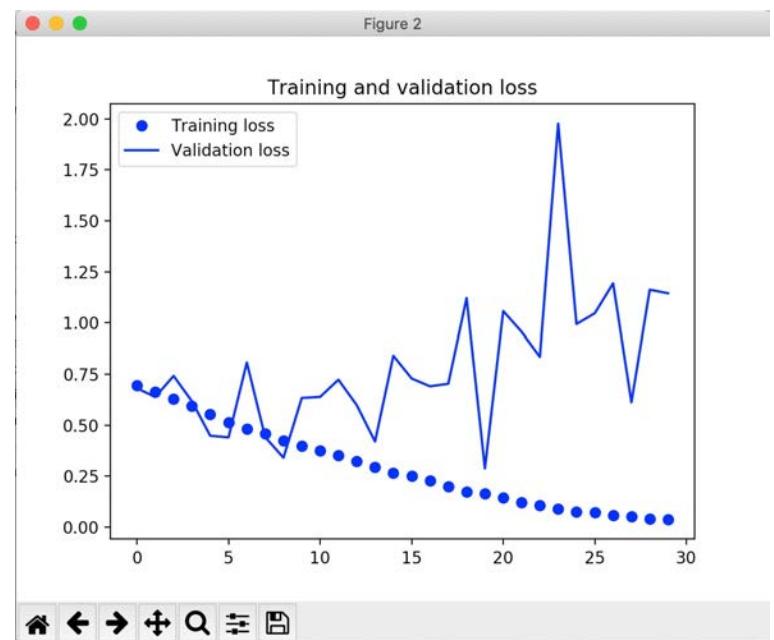
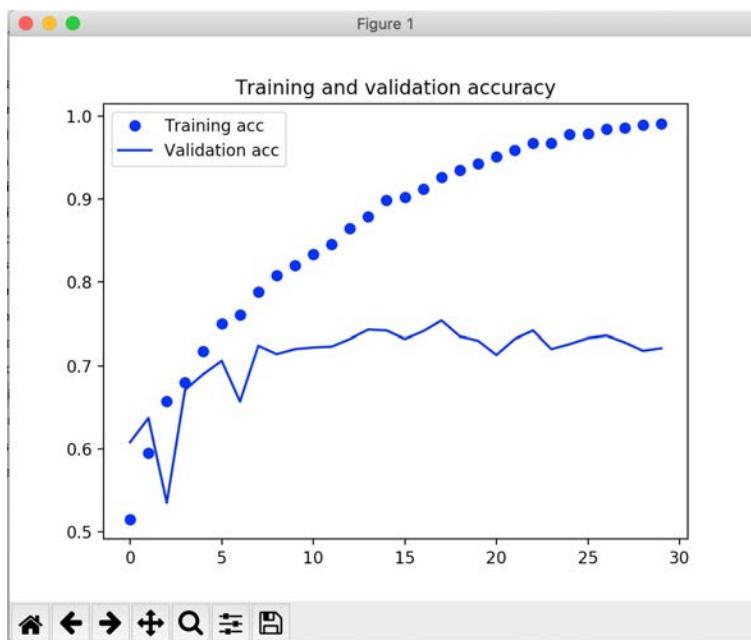
```
95
96 from tensorflow.keras import layers
97 from tensorflow.keras import models
98
99 model = models.Sequential()
100 model.add(layers.Conv2D(32, (3, 3),
101             activation='relu', input_shape=(150, 150,
102                 3)))
103 model.add(layers.MaxPooling2D((2, 2)))
104 model.add(layers.Conv2D(64, (3, 3),
105             activation='relu'))
106 model.add(layers.MaxPooling2D((2, 2)))
107 model.add(layers.Conv2D(128, (3, 3),
108             activation='relu'))
109 model.add(layers.MaxPooling2D((2, 2)))
110 model.add(layers.Flatten())
111 model.add(layers.Dense(512, activation='relu'))
112 model.add(layers.Dense(1, activation='sigmoid'))
113
114 print(model.summary())
115
116 from tensorflow.keras import optimizers
117
118 model.compile(loss='binary_crossentropy',
119                 optimizer=optimizers.RMSprop(lr=1e-4),
120                 metrics=['acc'])
121
122 # All images will be rescaled by 1./255
123 train_datagen =
124     ImageDataGenerator(rescale=1./255)
125 test_datagen = ImageDataGenerator(rescale=1./255)
126
127 train_generator =
128     train_datagen.flow_from_directory(
129         train_dir, target_size=(150, 150),
130         batch_size=20, class_mode='binary')
131 validation_generator =
132     test_datagen.flow_from_directory(
133         validation_dir, target_size=(150, 150),
134         batch_size=20, class_mode='binary')
135
136 for data_batch, labels_batch in train_generator:
137     print('data batch shape:', data_batch.shape)
138     print('labels batch shape:',
139         labels_batch.shape)
140     break
141
142 history = model.fit_generator(train_generator,
143                             steps_per_epoch=100, epochs=30,
144                             validation_data=validation_generator,
145                             validation_steps=50)
146
147 model
148     .save
149     ('/Users/swsun/Downloads/cats_and_dogs_small_
150     1.h5')
```

Full code (2/3)

```
136 import matplotlib.pyplot as plt
137
138 acc = history.history['acc']
139 val_acc = history.history['val_acc']
140 loss = history.history['loss']
141 val_loss = history.history['val_loss']
142
143 epochs = range(len(acc))
144
145 plt.plot(epochs, acc, 'bo', label='Training acc')
146 plt.plot(epochs, val_acc, 'b', label='Validation
147 acc')
148 plt.title('Training and validation accuracy')
149 plt.legend()
150
151 plt.figure()
152
153 plt.plot(epochs, loss, 'bo', label='Training
154 loss')
155 plt.plot(epochs, val_loss, 'b',
156 label='Validation loss')
157 plt.title('Training and validation loss')
158 plt.legend()
```

Practice 3

- Plot the results
 - Accuracy: training / validation
 - Loss: training / validation

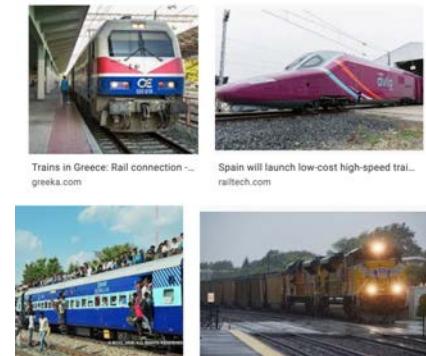


Homework 1

- Implementing a binary classification example
 - apples
 - bananas



- Implementing a multiclass classification example
 - cars, airplanes, trains, ...



Outline

- Introduction to convnets
- Training a convnet from scratch on a small dataset
- Using a pretrained convnet

Using a pretrained convnet

- According to a model
 - pretrained model previously
 - Pretrained from others
- **4 steps to predict the possible class**
 - Step 1: Using a trained model
 - Step 2: Load an image
 - Step 3: Convert the image
 - Step 4: Predict the class



It's a Dog!

Step 1: Using a trained model

- Load the convnet (CNN) model generated
 - previously



cats_and_dogs_small_2.h5

- Find out the **.h5** file (model) you generated

```
1 import os, shutil  
2 from tensorflow.keras.models import  
     load_model  
3 from tensorflow.keras.preprocessing.image  
     import load_img  
4 from tensorflow.keras.preprocessing.image  
     import img_to_array  
  
5  
6  
7 # load the model  
8 model = load_model('Users/swsun/Downloads/cats_and_dogs_  
     small_1.h5')  
     Load the model  
     assign the location  
     of the file
```

Step 2: Load an image

- Load the image file



```
10 #load the iamge          assign the name of the image file  
11 file_name='/Users/sunshih-wei/Documents/python/  
           dog.jpg'  
12 image = load_img(file_name,  
                   target_size=(150, 150))
```

Load the image

- Resize the image
 - to the proper size:
 - (150, 150): what you trained before

Step 3: Convert the image

- Convert the image file
 - To a proper tensor (array)

- Reshape the tensor
– as a 4D tensor $(1, 150, 150, 3)$

Step 4: Predict the class

- Using the model's method
 - model: an object
 - Predict which class it belongs

```
18 #predict the class: 0 is cat, 1 is dog  
19 print(model.predict(image))
```

Predict the class

Results:



StreamExecutor device (0): Host, Default Version

[[1]]

- Binary classification

- 2 classes: 0 or 1

- 0: cat

- 1: dog

- or inverted



StreamExecutor device (0): Host, Default Version

[[0]]

Full code

```
1 import os, shutil
2 from tensorflow.keras.models import
    load_model
3 from tensorflow.keras.preprocessing.image
    import load_img
4 from tensorflow.keras.preprocessing.image
    import img_to_array
5
6
7 # load the model
8 model =
    load_model
    ('/Users/swsun/Downloads/cats_and_dogs_
    small_1.h5')
9
10 #load the iamge
11 file_name='/Users/sunshih-wei/Documents/python/
    dog.jpg'
12 image = load_img(file_name,
    target_size=(150,150))
13
14 #image to array, reshape
15 image = img_to_array(image)
16 image = image.reshape((1, image.shape[0],
    image.shape[1], image.shape[2]))
17
18 #predict the class: 0 is cat, 1 is dog
19 print(model.predict(image))
```

Practice 1

- Predict the photos
 - By your own model
- Display: dog / cat
 - Hint:
 - if ... 0, else ...
- Download other photos



It's a dog!



It's a cat!

Textbook Reading

- *Deep Learning with Python*
 - Ch 8, Introduction to deep learning for computer vision
 - 8.1, Introduction to convnets
 - 8.2 Training a convnet from scratch on a small dataset
 - p. 201 - p. 224